

# Annotation of Calibration Patterns for RGB-LiDAR Evaluations using Segmentation Models

BRUNO SILVA<sup>1</sup>, PhD. Student

GONÇALO RIBEIRO<sup>1</sup>, PhD. Student

<sup>1</sup>Department of Mechanical Engineering, University of Aveiro

## ABSTRACT

Extrinsic calibration is a crucial step in sensor fusion. This work presents a deep-learning-based augmentation to ATOM, a multi-sensor and multi-modal calibration framework, for automating the annotation of the outer borders of calibration patterns, eliminating the need for manual labeling. Instead of predicting the outer corners directly with a CNN/FC combo, which proved infeasible due to variable output sizes, we trained segmentation models to detect the pattern and later extracted the edges using classical computer vision techniques. After testing various architectures, a U-Net model with a pretrained ResNet50 backbone delivered the best results.

## INDEX TERMS

Extrinsic Calibration, Robot Calibration, Calibration Pattern, Machine Learning

## I. INTRODUCTION

Extrinsic calibration is a fundamental process in robotics vision that involves determining the relative pose (position and orientation) between different sensors, known as *sensor to sensor calibration*, or between a sensor and a known reference frame, which is known as *sensor to coordinate frame*. Extrinsic calibration is crucial because it allows for the accurate integration of data from multiple sensors, enabling sensor fusion. For instance, in an autonomous vehicle, the visual information of the camera needs to be accurately aligned with the distance measurements of the LiDAR to build a coherent understanding of the surroundings. Similarly, in robot arms, the position of the camera position relative to the end-effector must be precisely known to perform tasks like object manipulation.

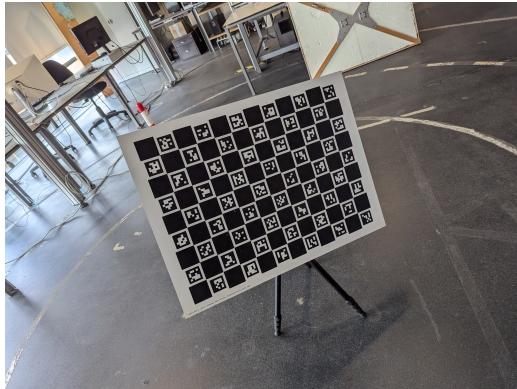
Usually, iterative approaches are used. These rely on a cost function specific to a sensor modality but usually suffer from ambiguities in some way or another. A typical cost function for an RGB camera relies on computing the difference between the projection of the detection in the 2D image of some key points into a coordinate frame where these key points are precisely known. The objects that contain these precisely known points are called *Calibration patterns*. The most common types are chessboards and *ChArUcos*. *ChArUcos* are chessboards with unique identifiable symbols embedded on each square that allow computer vision algorithms to decipher if the pattern is upside down or if the framing of the image cuts part of the calibration object off. The issue with the aforementioned RGB cost function is that they have multiple local minima and not always converge. A simple example is to picture an RGB camera fixed on the end of a prismatic joint with the calibration pattern in front of the sensor, perfectly

perpendicular to it. Both moving the offset of the joint or moving the RGB sensor on the mount can lead to the same relative distance between the sensor and the pattern, leading to equal detections of the key points, thus leading to ambiguity.

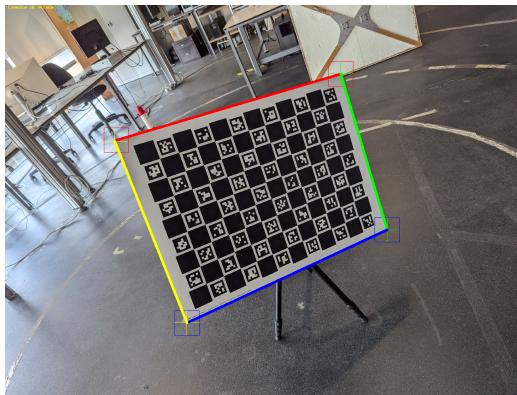
Despite the shortcomings mentioned, these cost functions have the advantage of generally performing more accurately in larger systems, as there is much more variety of data, and can still be used effectively in most simpler systems, as their requirements are only the existence of a sensor and a pattern. However, this creates the need for true evaluation procedures to assess the quality of the calibration results and to make them comparable in order to be publishable.

This project tackles an improvement to a previously fully manual and cumbersome evaluation method between RGB and LiDAR sensors, integrated on ATOM, a well established multi-sensor multi-modal calibration framework in the scientific community. The working principle is that by knowing the physical outer limits of the pattern in the 2D image, these points can be projected into the coordinate frame of LiDAR sensor, provided that the camera intrinsics are known. Afterward, the 3D points resulting from the projection can be directly compared with the 3D points of the outer border of the pattern. This comparison is not ambiguous as the projected points only line up in the 3D frame if the geometric transformations required for the projection are indeed correct. In ATOM, the 3D border points are already priorly labeled as they are a requirement for the cost function that optimizes the pose of the 3D LiDAR sensors. However, the 2D points of the border are not required by the RGB cost function. The current solution is a manual border labeling method, as a simple automatic method would struggle with detecting orientation and deal with edge cases. This project aims to

develop an automatic method using deep learning to simplify the evaluation pipeline to the user. Figure 1 and Figure 2 represent an example input and a desired output. On the output, each color encodes the border of one of the sides of the calibration pattern.



**FIGURE 1.** Input image



**FIGURE 2.** Desired Output, where red:top ; green:right; blue:bottom, green:right.

## II. PROPOSED APPROACH

### A. DATA PREPARATION

We initially had a dataset consisting of images with labeled outer corners of a ChArUco board, specifically the four corners that define the physical boundary of the pattern. To improve dataset diversity, we captured additional images featuring different types of calibration patterns.

To facilitate data annotation, we developed a program to label these outer corners bypassing the usual calibration pipeline necessary to obtain these annotations. Additionally, we created a program to convert the labeled corners into segmentation masks, which were later used for model training.

### B. WHY USE SEGMENTATION MODELS

Using a simpler CNN/FC combination to directly predict the coordinates of the outer corners of the ChArUco board was not feasible because the number of output points is not

constant. Designing a neural network with dynamic outputs would introduce significant complexity.

In many cases, at least one of the board's outer corners is clipped in the image. When this happens, five points are needed instead of four to properly define the pattern's edges. If two adjacent corners are clipped, only three sides remain visible.

To address this challenge, we opted to predict the segmentation mask of the pattern instead. This allowed us to later extract the edges using classical computer vision techniques.

### C. MODEL SELECTION AND TRAINING SETUP

To select the models to use, we prioritized the source code availability, as well as the availability of pre-trained weights. Currently the PyTorch Vision framework has three models for semantic segmentation: DeepLabV3 [1], FCN [2] and LRASPP [3]. From this models we choose DeepLabV3 since the authors claimed the best results. In addition to the PyTorch models, we also looked for a simple well known model in order to have an easier model to use. For that we choose the U-Net [unet].

### D. TRANSFER LEARNING WITH DEEPLABV3 AND RESNET50 BACKBONE

Using transfer learning with a pretrained DeepLabv3 model was the first option considered. DeepLabV3 is a complex model designed for image segmentation. Including the Resnet50 backbone, which is a CNN, the model has in total around 41M parameters. Upon freezing the backbone, we were left with 17M trainable parameters. The pretrained model came from PyTorch built-in models and was originally trained on a subset of known COCO dataset, using only the 20 categories that are present in the Pascal VOC dataset, mainly consisting of big every objects like cars, planes, sofas, bicycle...

The train was conducted using the Adam optimizer and cross entropy loss, batch size of 4 and 10 epochs. Our dataset had 80 images. We found the loss to stay constant and high, meaning our model was not learning anything from the training. We figured that the dataset that the model was originally trained on was too different from ours and the model could not learn within a feasible number of epochs with only 80 images. Our next attempt was to try a simpler segmentation model and train it from the ground up.

### E. U-NET TRAINED FROM SCRATCH

Training the previous model from scratch would require a lot of compute resources and probably be overkill for the desired task. As such, we decided to try a simpler model and train it from scratch. Unet is blablabla

Good results

### F. U-NET WITH PRETRAINED RESNET50 BACKBONE

Better results

---

### III. CONCLUSION

We presented a approach to automate the annotation of the outer borders of a calibration patterns using deep learning. By leveraging segmentation models instead of direct CNN/FC combos, we addressed challenges related to pattern occlusions and variable output sizes. While transfer learning with DeepLabV3 did not yield satisfactory results, training a U-Net model with a ResNet50 backbone proved highly effective. This approach streamlines the evaluation process, reducing manual effort while maintaining accuracy. Future work could focus on further refining the model.

### IV. REFERENCES

- [1] L.-C. Chen, G. Papandreou, F. Schroff, and H. Adam, *Rethinking atrous convolution for semantic image segmentation*, Dec. 5, 2017. DOI: 10.48550/arXiv.1706.05587. arXiv: 1706.05587[cs].
- [2] E. Shelhamer, J. Long, and T. Darrell, “Fully convolutional networks for semantic segmentation,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 39, no. 4, pp. 640–651, Apr. 2017. DOI: 10.1109/TPAMI.2016.2572683.
- [3] A. Howard *et al.*, “Searching for MobileNetV3,” in *2019 IEEE/CVF International Conference on Computer Vision (ICCV)*, Oct. 2019, pp. 1314–1324. DOI: 10.1109/ICCV.2019.00140.

• • •